

Analyze_ab_test_results_notebook

October 14, 2020

0.1 Analyze A/B Test Results

You may either submit your notebook through the workspace here, or you may work from your local machine and submit through the next page. Either way assure that your code passes the project [RUBRIC](#). **Please save regularly.**

This project will assure you have mastered the subjects covered in the statistics lessons. The hope is to have this project be as comprehensive of these topics as possible. Good luck!

0.2 Table of Contents

- Section ??
- Section ??
- Section ??
- Section ??

Introduction

A/B tests are very commonly performed by data analysts and data scientists. It is important that you get some practice working with the difficulties of these

For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

As you work through this notebook, follow along in the classroom and answer the corresponding quiz questions associated with each question. The labels for each classroom concept are provided for each question. This will assure you are on the right track as you work through the project, and you can feel more confident in your final submission meeting the criteria. As a final check, assure you meet all the criteria on the [RUBRIC](#).

Part I - Probability

To get started, let's import our libraries.

```
In [2]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

1. Now, read in the `ab_data.csv` data. Store it in `df`. Use your dataframe to answer the questions in Quiz 1 of the classroom.

a. Read in the dataset and take a look at the top few rows here:

```
In [3]: df = pd.read_csv('ab_data.csv')
df.head
```

```
Out[3]: <bound method NDFrame.head of
0      851104  2017-01-21 22:11:48.556739  control  old_page  0
1      804228  2017-01-12 08:01:45.159739  control  old_page  0
2      661590  2017-01-11 16:55:06.154213  treatment new_page  0
3      853541  2017-01-08 18:28:03.143765  treatment new_page  0
4      864975  2017-01-21 01:52:26.210827  control  old_page  1
5      936923  2017-01-10 15:20:49.083499  control  old_page  0
6      679687  2017-01-19 03:26:46.940749  treatment new_page  1
7      719014  2017-01-17 01:48:29.539573  control  old_page  0
8      817355  2017-01-04 17:58:08.979471  treatment new_page  1
9      839785  2017-01-15 18:11:06.610965  treatment new_page  1
10     929503  2017-01-18 05:37:11.527370  treatment new_page  0
11     834487  2017-01-21 22:37:47.774891  treatment new_page  0
12     803683  2017-01-09 06:05:16.222706  treatment new_page  0
13     944475  2017-01-22 01:31:09.573836  treatment new_page  0
14     718956  2017-01-22 11:45:11.327945  treatment new_page  0
15     644214  2017-01-22 02:05:21.719434  control  old_page  1
16     847721  2017-01-17 14:01:00.090575  control  old_page  0
17     888545  2017-01-08 06:37:26.332945  treatment new_page  1
18     650559  2017-01-24 11:55:51.084801  control  old_page  0
19     935734  2017-01-17 20:33:37.428378  control  old_page  0
20     740805  2017-01-12 18:59:45.453277  treatment new_page  0
21     759875  2017-01-09 16:11:58.806110  treatment new_page  0
22     767017  2017-01-12 22:58:14.991443  control  new_page  0
23     793849  2017-01-23 22:36:10.742811  treatment new_page  0
24     905617  2017-01-20 14:12:19.345499  treatment new_page  0
25     746742  2017-01-23 11:38:29.592148  control  old_page  0
26     892356  2017-01-05 09:35:14.904865  treatment new_page  1
27     773302  2017-01-12 08:29:49.810594  treatment new_page  0
28     913579  2017-01-24 09:11:39.164256  control  old_page  1
29     736159  2017-01-06 01:50:21.318242  treatment new_page  0
...     ...
294448  776137  2017-01-12 05:53:12.386730  treatment new_page  0
294449  883344  2017-01-22 23:15:58.645325  treatment new_page  0
294450  825594  2017-01-06 12:37:08.897784  treatment new_page  0
294451  875688  2017-01-14 07:19:49.042869  control  old_page  0
294452  927527  2017-01-12 10:52:11.084740  control  old_page  0
294453  789177  2017-01-17 18:17:56.215378  control  old_page  0
294454  937338  2017-01-19 03:23:22.236666  treatment new_page  0
294455  733101  2017-01-23 12:52:58.711914  treatment new_page  0
```

294456	679096	2017-01-02	16:43:49.237940	treatment	new_page	0
294457	691699	2017-01-09	23:42:35.963486	treatment	new_page	0
294458	807595	2017-01-22	10:43:09.285426	treatment	new_page	0
294459	924816	2017-01-20	10:59:03.481635	control	old_page	0
294460	846225	2017-01-16	15:24:46.705903	treatment	new_page	0
294461	740310	2017-01-10	17:22:19.762612	control	old_page	0
294462	677163	2017-01-03	19:41:51.902148	treatment	new_page	0
294463	832080	2017-01-19	13:18:27.352570	control	old_page	0
294464	834362	2017-01-17	01:51:56.106436	control	old_page	0
294465	925675	2017-01-07	20:38:26.346410	treatment	new_page	0
294466	923948	2017-01-09	16:33:41.104573	control	old_page	0
294467	857744	2017-01-05	08:00:56.024226	control	old_page	0
294468	643562	2017-01-02	19:20:05.460595	treatment	new_page	0
294469	755438	2017-01-18	17:35:06.149568	control	old_page	0
294470	908354	2017-01-11	02:42:21.195145	control	old_page	0
294471	718310	2017-01-21	22:44:20.378320	control	old_page	0
294472	822004	2017-01-04	03:36:46.071379	treatment	new_page	0
294473	751197	2017-01-03	22:28:38.630509	control	old_page	0
294474	945152	2017-01-12	00:51:57.078372	control	old_page	0
294475	734608	2017-01-22	11:45:03.439544	control	old_page	0
294476	697314	2017-01-15	01:20:28.957438	control	old_page	0
294477	715931	2017-01-16	12:40:24.467417	treatment	new_page	0

[294478 rows x 5 columns]>

b. Use the cell below to find the number of rows in the dataset.

```
In [5]: df.shape[0]
```

```
Out[5]: 294478
```

c. The number of unique users in the dataset.

```
In [7]: df.nunique()
```

```
Out[7]: user_id      290584
        timestamp    294478
        group         2
        landing_page   2
        converted       2
        dtype: int64
```

d. The proportion of users converted.

```
In [8]: size = df[df['converted'] == 1].shape[0]
        print(size)
```

```
35237
```

```
In [10]: proportion = (size / df.shape[0])*100
         print(proportion)
```

```
11.96591935560551
```

e. The number of times the `new_page` and `treatment` don't match.

```
In [12]: df.query('group == "treatment" and landing_page != "new_page").shape[0] + \
         df.query('group == "control" and landing_page != "old_page").shape[0]
```

```
Out[12]: 3893
```

f. Do any of the rows have missing values?

```
In [13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
user_id      294478 non-null int64
timestamp    294478 non-null object
group        294478 non-null object
landing_page 294478 non-null object
converted    294478 non-null int64
dtypes: int64(2), object(3)
memory usage: 11.2+ MB
```

2. For the rows where **treatment** does not match with **new_page** or **control** does not match with **old_page**, we cannot be sure if this row truly received the new or old page. Use **Quiz 2** in the classroom to figure out how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in **df2**.

```
In [18]: df2 = df.query("group == 'control' and landing_page == 'old_page'")
         df2 = df2.append(df.query("group == 'treatment' and landing_page == 'new_page'))
```

```
In [19]: # Double Check all of the correct rows were removed - this should be 0
         df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) == False].shape[0]
```

```
Out[19]: 0
```

3. Use **df2** and the cells below to answer questions for **Quiz3** in the classroom.

a. How many unique **user_ids** are in **df2**?

```
In [20]: df2.user_id.nunique()
```

```
Out[20]: 290584
```

b. There is one **user_id** repeated in **df2**. What is it?

```
In [21]: df2[df2['user_id'].duplicated()]['user_id']
```

```
Out[21]: 2893      773192
         Name: user_id, dtype: int64
```

c. What is the row information for the repeat **user_id**?

```
In [22]: df2[df2['user_id'].duplicated()]
```

```
Out[22]:
```

	user_id	timestamp	group	landing_page	converted	
	2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

d. Remove **one** of the rows with a duplicate **user_id**, but keep your dataframe as **df2**.

```
In [23]: df2 = df2.drop(2893)
```

4. Use **df2** in the cells below to answer the quiz questions related to **Quiz 4** in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

```
In [24]: df2.converted.mean()
```

```
Out[24]: 0.11965919355605512
```

b. Given that an individual was in the control group, what is the probability they converted?

```
In [26]: control_prop = df2.query("group == 'control'")['converted'].mean()
         control_prop
```

```
Out[26]: 0.1203863045004612
```

c. Given that an individual was in the treatment group, what is the probability they converted?

```
In [27]: treatment_prop = df2.query("group == 'treatment'")['converted'].mean()
         treatment_prop
```

```
Out[27]: 0.11880806551510564
```

d. What is the probability that an individual received the new page?

```
In [28]: df2.query('landing_page == "new_page"').shape[0] / df2.shape[0]
```

```
Out[28]: 0.5000619442226688
```

e. Consider your results from parts (a) through (d) above, and explain below whether you think there is sufficient evidence to conclude that the new treatment page leads to more conversions.

No, there is no sufficient evidence to say that the new treatment page leads to more conversions.

Because the probability that users in group of treatment will convert is actually slightly less than the probability that user in the group of control will convert.

Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your hypothesis in terms of words or in terms of p_{old} and p_{new} , which are the converted rates for the old and new pages.

Null hypotheses : * $H_0 : p_{old} \geq p_{new}$ * $H_1 : p_{old} < p_{new}$

in other words : * $H_0 : p_{new} \leq p_{old}$ * $H_1 : p_{new} > p_{old}$

2. Assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page.

Use a sample size for each page equal to the ones in **ab_data.csv**.

Perform the sampling distribution for the difference in **converted** between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn't make complete sense right now, don't worry - you are going to work through the problems below to complete this problem. You can use **Quiz 5** in the classroom to make sure you are on the right track.

a. What is the **conversion rate** for p_{new} under the null?

```
In [31]: p_new = df2['converted'].mean()
         print(p_new)
```

```
0.119597087245
```

b. What is the **conversion rate** for p_{old} under the null?

```
In [33]: p_null
```

```
Out[33]: 0.11959708724499628
```

c. What is n_{new} , the number of individuals in the treatment group?

```
In [34]: n_new = df2.query("landing_page == 'new_page'").shape[0]
         print(n_new)
```

```
145310
```

d. What is n_{old} , the number of individuals in the control group?

```
In [35]: n_old = df2.query("landing_page == 'old_page').shape[0]
         print(n_old)
```

```
145274
```

e. Simulate n_{new} transactions with a conversion rate of p_{new} under the null. Store these n_{new} 1's and 0's in **new_page_converted**.

```
In [ ]: new_page_converted = np.random.binomial(1, p_null, n_new)
```

f. Simulate n_{old} transactions with a conversion rate of p_{old} under the null. Store these n_{old} 1's and 0's in **old_page_converted**.

```
In [37]: old_page_converted = np.random.binomial(1, p_null, n_old)
```

g. Find $p_{new} - p_{old}$ for your simulated values from part (e) and (f).

```
In [47]: old_page_converted.mean() - new_page_converted.mean()
```


```
Out[47]: -0.00040369307873495963
```

h. Create 10,000 $p_{new} - p_{old}$ values using the same simulation process you used in parts (a) through (g) above. Store all 10,000 values in a NumPy array called **p_diffs**.

```
In [52]: p_diffs = []
         new_converted_simulation = np.random.binomial(n_new, p_null, 10000)/n_new
         old_converted_simulation = np.random.binomial(n_old, p_null, 10000)/n_old
         p_diffs = new_converted_simulation - old_converted_simulation
```

i. Plot a histogram of the **p_diffs**. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

```
In [53]: plt.hist(p_diffs);
```



- j. What proportion of the **p_diffs** are greater than the actual difference observed in **ab_data.csv**?

```
In [55]: act_diff = df[df['group'] == 'treatment']['converted'].mean() - df[df['group'] == 'control']['converted'].mean()
          print(act_diff)
```

```
-0.00147959979408
```

```
In [56]: p_diffs = np.array(p_diffs)
          p_diffs
```

```
Out[56]: array([ 7.13626516e-04, -6.55786080e-04, -7.76860379e-05, ...,
                  7.13740776e-04, -5.80171122e-04,  1.31940740e-03])
```

```
In [57]: (act_diff < p_diffs).mean()
```

```
Out[57]: 0.8921999999999999
```

- k. Please explain using the vocabulary you've learned in this course what you just computed in part j. What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

I've computed the p-value in j, which is the probability that we will observe this statistic, given the null hypothesis is true.

Since the p-value is large here, So we fail to reject the null hypothesis.

- l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer to the number of rows associated with the old page and new pages, respectively.

```
In [58]: import statsmodels.api as sm
```

```
convert_old = df2.query('group == "control"')['converted'].sum()
convert_new = df2.query('group == "treatment"')['converted'].sum()
n_old = df2.query('landing_page == "old_page"').shape[0]
n_new = df2.query('landing_page == "new_page"').shape[0]
print(convert_old, convert_new, n_old, n_new)
```

```
/opt/conda/lib/python3.6/site-packages/statsmodels/compat/pandas.py:56: FutureWarning: The pandas
from pandas.core import datetools
```

```
17489 17264 145274 145310
```

- m. Now use `stats.proportions_ztest` to compute your test statistic and p-value. [Here](#) is a helpful link on using the built in.

```
In [97]: z_score, p_value = sm.stats.proportions_ztest(np.array([convert_new, convert_old])
                                                    ,np.array([n_new, n_old]), alternative =
                                                    z_score, p_value
```

```
Out[97]: (-1.3109241984234394, 0.90505831275902449)
```

- n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts j. and k.?

The z-score here is -1.31 inside our critical value of 1.959 and the p-value is still large, So it is likely that our statistic is from the null

This means the z-score and p-value agree with the findings in parts j and k that we cannot reject the null hypothesis.

Part III - A regression approach

1. In this final part, you will see that the result you achieved in the A/B test in Part II above can also be achieved by performing regression.

- a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

Logistic regression.

- b. The goal is to use **statsmodels** to fit the regression model you specified in part a. to see if there is a significant difference in conversion based on which page a customer receives. However, you first need to create in `df2` a column for the intercept, and create a dummy variable column for which page each user received. Add an **intercept** column, as well as an **ab_page** column, which is 1 when an individual receives the **treatment** and 0 if **control**.

```
In [189]: df2['intercept']=1
          df2['ab_page']=0
          ab_page_index = df2[df2['group']=='treatment'].index
          df2.loc[ab_page_index, "ab_page"] = 1
```

```
df2.head()
```

```
Out[189]:
```

	timestamp	group	landing_page	converted	ab_page \
user_id					
851104	2017-01-21 22:11:48.556739	control	old_page	0	0
804228	2017-01-12 08:01:45.159739	control	old_page	0	0
864975	2017-01-21 01:52:26.210827	control	old_page	1	0
936923	2017-01-10 15:20:49.083499	control	old_page	0	0
719014	2017-01-17 01:48:29.539573	control	old_page	0	0

	intercept	country	us	uk	ca	US	UK	CA
user_id								
851104	1	US	0	0	1	1	0	0
804228	1	US	0	0	1	1	0	0
864975	1	US	0	0	1	1	0	0
936923	1	US	0	0	1	1	0	0
719014	1	US	0	0	1	1	0	0

- c. Use **statsmodels** to instantiate your regression model on the two columns you created in part b., then fit the model using the two columns you created in part b. to predict whether or not an individual converts.

```
In [190]: lm = sm.OLS(df2['converted'], df2[['intercept', 'ab_page']])
          results=lm.fit()
          results.summary()
```

```
Out[190]: <class 'statsmodels.iolib.summary.Summary'>
        """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  converted    R-squared:                  0.000
Model:                            OLS      Adj. R-squared:              0.000
Method:                 Least Squares    F-statistic:                  1.719
Date:                Tue, 13 Oct 2020    Prob (F-statistic):          0.190
Time:                  18:43:51          Log-Likelihood:              -85267.
No. Observations:                290584    AIC:                        1.705e+05
Df Residuals:                    290582    BIC:                        1.706e+05
Df Model:                            1
```

```

Covariance Type:      nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
intercept      0.1204      0.001     141.407      0.000      0.119      0.122
ab_page     -0.0016      0.001     -1.311      0.190     -0.004      0.001
=====
Omnibus:            125553.456   Durbin-Watson:           2.000
Prob(Omnibus):            0.000   Jarque-Bera (JB):        414313.355
Skew:              2.345   Prob(JB):              0.00
Kurtosis:           6.497   Cond. No.              2.62
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly speci
"""

```

- d. Provide the summary of your model below, and use it as necessary to answer the following questions.
- e. What is the p-value associated with **ab_page**? Why does it differ from the value you found in **Part II**? **Hint:** What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in **Part II**?

Our hypothesis here is:

$$H_0 : p_{new} - p_{old} = 0$$

$$H_1 : p_{new} - p_{old} \neq 0$$

- f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

We should consider other factors into the regression model as they might influence the conversions too. For instance student segments [new v/s returning candidates] might create change aversion or even, the opposite as a predisposition to conversion. Seasonality like new terms or New years might mean more interest in new skills/ resolutions. Timestamps are included but without regionality, they do not indicate if seasonality was a factor or not. [as different countries follow different term and weather patterns. Factors like device on which tests were taken or course which was looked at, prior academic background, age, might alter experience and ultimately, conversions. These are limitations which should be at least kept in mind while making the final decision. The disadvantages to adding additional terms into the regression model is that even with additional factors we can never account for all influencing factors or accomodate them. Plus, small pilots and pivots sometimes work better in practice than long-drawn research without execution.

- g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives in. You will need to read in the **countries.csv** dataset and merge together your datasets on the appropriate rows. [Here](#) are the docs for joining tables.

Does it appear that country had an impact on conversion? Don't forget to create dummy variables for these country columns - **Hint: You will need two columns for the three dummy variables.** Provide the statistical output as well as a written response to answer this question.

```
In [173]: df_countries = pd.read_csv('countries.csv')
          df_countries.head()
```

```
Out[173]:
```

	user_id	country
0	834778	UK
1	928468	US
2	822059	UK
3	711597	UK
4	710616	UK

```
In [175]: df2[['CA', 'UK', 'US']] = pd.get_dummies(df2['country'])
          df2
```

```
Out[175]:
```

	user_id	timestamp	group	landing_page	converted	\
	851104	2017-01-21 22:11:48.556739	control	old_page	0	
	804228	2017-01-12 08:01:45.159739	control	old_page	0	
	864975	2017-01-21 01:52:26.210827	control	old_page	1	
	936923	2017-01-10 15:20:49.083499	control	old_page	0	
	719014	2017-01-17 01:48:29.539573	control	old_page	0	
	644214	2017-01-22 02:05:21.719434	control	old_page	1	
	847721	2017-01-17 14:01:00.090575	control	old_page	0	
	650559	2017-01-24 11:55:51.084801	control	old_page	0	
	935734	2017-01-17 20:33:37.428378	control	old_page	0	
	746742	2017-01-23 11:38:29.592148	control	old_page	0	
	913579	2017-01-24 09:11:39.164256	control	old_page	1	
	690284	2017-01-13 17:22:57.182769	control	old_page	0	
	710349	2017-01-11 22:24:44.226492	control	old_page	0	
	677533	2017-01-23 17:48:50.491821	control	old_page	0	
	831737	2017-01-11 21:18:20.911015	control	old_page	1	
	771087	2017-01-16 00:05:29.983919	control	old_page	0	
	896163	2017-01-22 09:10:20.753218	control	old_page	0	
	862225	2017-01-08 14:49:37.335432	control	old_page	1	
	939593	2017-01-05 09:15:31.984283	control	old_page	0	
	702260	2017-01-18 13:55:31.488221	control	old_page	0	
	670941	2017-01-05 08:16:41.306478	control	old_page	0	
	850231	2017-01-18 17:18:04.790584	control	old_page	1	
	685794	2017-01-20 14:54:58.150621	control	old_page	0	
	714733	2017-01-03 08:22:37.904146	control	old_page	0	
	710967	2017-01-10 02:19:22.842142	control	old_page	0	
	680201	2017-01-11 10:38:45.952887	control	old_page	0	
	790863	2017-01-19 11:02:39.220320	control	old_page	0	
	717595	2017-01-23 18:19:08.148166	control	old_page	0	
	779854	2017-01-11 21:28:30.735359	control	old_page	0	
	916307	2017-01-19 17:27:38.676600	control	old_page	0	

...
924332	2017-01-15 19:38:52.858024	treatment	new_page	0
849625	2017-01-06 17:54:07.563311	treatment	new_page	0
929723	2017-01-10 15:13:48.352399	treatment	new_page	0
774769	2017-01-03 06:01:36.251836	treatment	new_page	0
733871	2017-01-21 17:54:08.810964	treatment	new_page	1
844588	2017-01-16 20:48:19.567178	treatment	new_page	0
641244	2017-01-07 16:57:26.193171	treatment	new_page	0
676072	2017-01-14 17:26:02.495442	treatment	new_page	0
886374	2017-01-07 13:43:39.202634	treatment	new_page	0
676732	2017-01-03 23:06:45.459467	treatment	new_page	0
862218	2017-01-04 10:43:07.846494	treatment	new_page	0
798826	2017-01-23 16:50:13.788528	treatment	new_page	0
836721	2017-01-12 17:37:50.966955	treatment	new_page	0
844901	2017-01-15 17:46:36.622726	treatment	new_page	0
653124	2017-01-14 13:44:51.745491	treatment	new_page	0
909437	2017-01-18 14:49:49.064452	treatment	new_page	0
776137	2017-01-12 05:53:12.386730	treatment	new_page	0
883344	2017-01-22 23:15:58.645325	treatment	new_page	0
825594	2017-01-06 12:37:08.897784	treatment	new_page	0
937338	2017-01-19 03:23:22.236666	treatment	new_page	0
733101	2017-01-23 12:52:58.711914	treatment	new_page	0
679096	2017-01-02 16:43:49.237940	treatment	new_page	0
691699	2017-01-09 23:42:35.963486	treatment	new_page	0
807595	2017-01-22 10:43:09.285426	treatment	new_page	0
846225	2017-01-16 15:24:46.705903	treatment	new_page	0
677163	2017-01-03 19:41:51.902148	treatment	new_page	0
925675	2017-01-07 20:38:26.346410	treatment	new_page	0
643562	2017-01-02 19:20:05.460595	treatment	new_page	0
822004	2017-01-04 03:36:46.071379	treatment	new_page	0
715931	2017-01-16 12:40:24.467417	treatment	new_page	0

	ab_page	intercept	country	us	uk	ca	US	UK	CA
user_id									
851104	0	1	US	0	0	1	1	0	0
804228	0	1	US	0	0	1	1	0	0
864975	0	1	US	0	0	1	1	0	0
936923	0	1	US	0	0	1	1	0	0
719014	0	1	US	0	0	1	1	0	0
644214	0	1	US	0	0	1	1	0	0
847721	0	1	US	0	0	1	1	0	0
650559	0	1	CA	1	0	0	0	0	1
935734	0	1	US	0	0	1	1	0	0
746742	0	1	US	0	0	1	1	0	0
913579	0	1	US	0	0	1	1	0	0
690284	0	1	US	0	0	1	1	0	0
710349	0	1	UK	0	1	0	0	1	0
677533	0	1	US	0	0	1	1	0	0

831737	0	1	UK	0	1	0	0	1	0
771087	0	1	US	0	0	1	1	0	0
896163	0	1	UK	0	1	0	0	1	0
862225	0	1	US	0	0	1	1	0	0
939593	0	1	US	0	0	1	1	0	0
702260	0	1	US	0	0	1	1	0	0
670941	0	1	US	0	0	1	1	0	0
850231	0	1	US	0	0	1	1	0	0
685794	0	1	US	0	0	1	1	0	0
714733	0	1	US	0	0	1	1	0	0
710967	0	1	US	0	0	1	1	0	0
680201	0	1	US	0	0	1	1	0	0
790863	0	1	US	0	0	1	1	0	0
717595	0	1	US	0	0	1	1	0	0
779854	0	1	US	0	0	1	1	0	0
916307	0	1	UK	0	1	0	0	1	0
...
924332	1	1	US	0	0	1	1	0	0
849625	1	1	US	0	0	1	1	0	0
929723	1	1	CA	1	0	0	0	0	1
774769	1	1	US	0	0	1	1	0	0
733871	1	1	US	0	0	1	1	0	0
844588	1	1	US	0	0	1	1	0	0
641244	1	1	US	0	0	1	1	0	0
676072	1	1	US	0	0	1	1	0	0
886374	1	1	US	0	0	1	1	0	0
676732	1	1	UK	0	1	0	0	1	0
862218	1	1	US	0	0	1	1	0	0
798826	1	1	US	0	0	1	1	0	0
836721	1	1	US	0	0	1	1	0	0
844901	1	1	US	0	0	1	1	0	0
653124	1	1	CA	1	0	0	0	0	1
909437	1	1	UK	0	1	0	0	1	0
776137	1	1	US	0	0	1	1	0	0
883344	1	1	CA	1	0	0	0	0	1
825594	1	1	UK	0	1	0	0	1	0
937338	1	1	UK	0	1	0	0	1	0
733101	1	1	US	0	0	1	1	0	0
679096	1	1	US	0	0	1	1	0	0
691699	1	1	US	0	0	1	1	0	0
807595	1	1	US	0	0	1	1	0	0
846225	1	1	US	0	0	1	1	0	0
677163	1	1	US	0	0	1	1	0	0
925675	1	1	US	0	0	1	1	0	0
643562	1	1	CA	1	0	0	0	0	1
822004	1	1	CA	1	0	0	0	0	1
715931	1	1	UK	0	1	0	0	1	0

```

[290584 rows x 13 columns]

In [176]: df_countries.country.unique()

Out[176]: array(['UK', 'US', 'CA'], dtype=object)

In [178]: country_dummies = pd.get_dummies(df_countries['country'])
df_new = df_countries.join(country_dummies)

In [179]: df_new.head()

Out[179]:
   user_id country  CA  UK  US
0   834778      UK   0   1   0
1   928468      US   0   0   1
2   822059      UK   0   1   0
3   711597      UK   0   1   0
4   710616      UK   0   1   0

In [182]: lm = sm.OLS(df2['converted'], df2[['intercept', 'UK', 'US']])
results = lm.fit()
results.summary()

Out[182]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
Dep. Variable:                  converted    R-squared:                0.000
Model:                            OLS      Adj. R-squared:            0.000
Method:                 Least Squares    F-statistic:                 1.605
Date:                Tue, 13 Oct 2020    Prob (F-statistic):          0.201
Time:                  18:38:07          Log-Likelihood:            -85267.
No. Observations:                290584    AIC:                        1.705e+05
Df Residuals:                    290581    BIC:                        1.706e+05
Df Model:                          2
Covariance Type:                nonrobust
=====
                                coef    std err          t      P>|t|      [0.025     0.975]
-----
intercept          0.1153      0.003    42.792     0.000      0.110      0.121
UK                  0.0053      0.003     1.787     0.074     -0.001      0.011
US                  0.0042      0.003     1.516     0.130     -0.001      0.010
=====
Omnibus:                 125552.384    Durbin-Watson:                2.000
Prob(Omnibus):              0.000    Jarque-Bera (JB):             414306.036
Skew:                      2.345    Prob(JB):                      0.00
Kurtosis:                  6.497    Cond. No.                      9.94
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly speci
"""

```

- h. Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there significant effects on conversion. Create the necessary additional columns, and fit the new model.

Provide the summary results, and your conclusions based on the results.

Finishing Up

Congratulations! You have reached the end of the A/B Test Results project! You should be very proud of all you have accomplished!

Tip: Once you are satisfied with your work here, check over your report to make sure that it satisfies all the areas of the rubric (found on the project submission page at the end of the lesson). You should also probably remove all of the "Tips" like this one so that the presentation is as polished as possible.

0.3 Directions to Submit

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File > Download as** sub-menu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you've done this, you can submit your project by clicking on the "Submit Project" button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

```
In [ ]: from subprocess import call
        call(['python', '-m', 'nbconvert', 'Analyze_ab_test_results_notebook.ipynb'])
```